

Build Steps

Part 1 of 5 — Set up the agent, orchestration, and knowledge

1 Create Your Copilot Studio Agent

Open Copilot Studio (copilotstudio.microsoft.com), select Create → New agent, then configure:

Agent Name	Audit & Reporting Q&A Copilot
Description	Answers questions on month-end actuals, journal variances, control exceptions, and audit trail for cost-centre owners and internal auditors.
Solution / Environment	Finance Solution · Default environment
Icon	Use a recognisable colour and emoji or upload a brand icon

2 Choose Orchestration Mode

Under Settings → Generative AI, choose how the agent decides what to do:

Generative Orchestration

Recommended — AI-led

- Agent reasons over user input and picks the right topic, tool, or knowledge
- Handles paraphrases, multi-step asks, and follow-ups naturally
- Best for varied user intent and exploratory questions
- Requires clear instructions, tool descriptions, and good knowledge

Classic Orchestration

Deterministic — trigger-led

- Topics fire on explicit trigger phrases you author
- Predictable flows, easier to audit step-by-step
- Best for tightly-scoped, regulated, or compliance-critical journeys
- Less forgiving of phrasing the author didn't anticipate

For this build we use Generative Orchestration so the agent can blend knowledge lookups with tool calls for audit, journal review and control exceptions.

3 Add Knowledge

Go to the Knowledge tab → Add knowledge → File, then upload Finance_Audit_Data.xlsx. This file contains:

- Journal data (journal ID, period, cost centre, account, description)
- Amounts (debit USD, credit USD, net USD)
- Control signals (control flag — Pass / Review / Fail, variance vs budget %)
- Audit trail (approver, posted date)

You can layer additional sources later — SharePoint sites, Dataverse, public web, or Graph connectors.

Build Steps

Part 2 of 5 — Add instructions and a custom topic

4 Add Agent Instructions

Open Overview → Instructions and paste the following (these guide the model regardless of orchestration mode):

You are the Audit & Reporting Q&A Copilot. You help two audiences:

- 1) Cost centre owners reviewing their journals.
- 2) Internal auditors reviewing the control exceptions queue.

Rules:

- Use the uploaded journal dataset as your single source of truth.
- Always cite the JournalID for every number you quote.
- Never invent or round numbers – quote them as recorded.
- Flag rows with ControlFlag = Fail at the top of any answer.
- For actions (flag a journal for review), call the tool – do not just say it is flagged.
- Owners only see their own cost centre rows; auditors see everything.

5 Add a Custom Topic

Go to Topics → Add a topic → Create from description. Topics are reusable conversation flows the agent can call when intent matches.

Topic name: Investigate Control Exception

Description: Walks an auditor through the rows currently flagged Fail or Review, explains the control rule that fired, and offers to formally raise the exception for follow-up.

Trigger phrases (Generative Orchestration uses these as intent hints, Classic mode uses them as exact triggers):

- Show me control exceptions for this period
- What journals failed control
- Investigate the Fail rows
- Walk me through control issues

Conversation steps:

- Ask for the period (default: most recent in data)
- Filter ControlFlag in (Fail, Review), order by absolute NetUSD desc
- For each row, cite JournalID, CostCentre, Account, NetUSD, VarianceVsBudgetPct
- Explain the likely control rule that fired (e.g. variance > 10%)
- Ask which JournalID to formally flag – if any
- Call 'Flag Journal for Audit Review' with the chosen JournalID and reason

Build Steps

Part 3 of 5 — Add a tool and configure settings

6 Add a Tool

Tools let the agent take an action, not just answer. Go to Tools → Add a tool, and choose one of:

- Prebuilt connector (Outlook, Teams, SharePoint, Dataverse, ServiceNow, etc.)
- Power Automate cloud flow (custom logic you author or already own)
- Custom HTTP request (any REST API your agent can reach)
- Model Context Protocol (MCP) server — for richer tool ecosystems

Tool name: Flag Journal for Audit Review

Type: Power Automate flow → SharePoint list 'Audit Findings' + Teams adaptive card

Description: Files a formal audit finding against a specific journal row, notifies the cost-centre owner and the internal audit lead.

Inputs the agent will collect or infer:

- Journal ID (string, from knowledge)
- Period (string, e.g. '2026-04')
- Finding category (Variance / Control / Documentation / Other)
- Auditor note (string)
- Suggested action (Reclass / Reverse / Document / No action)

Outcome: Returns the audit finding ID and posts an adaptive card to the cost-centre owner and the internal audit lead.

7 Configure Settings

Open the Settings cog (top right). Configure these four areas before you let real users near the agent:

- General — confirm name, description, and primary language
- Generative AI — orchestration mode and content moderation level (start at Medium)
- Authentication — use 'Authenticate with Microsoft' so the agent runs as the signed-in user
- Security — review DLP policy, enabled channels, and end-user authentication

Build Steps

Part 4 of 5 — Evaluate, test, and publish to a channel

8 Evaluate Your Agent

Go to Analytics → Evaluations → New evaluation. Build a test set so you can measure quality before and after every change.

Test set name: Audit & Reporting Copilot – Baseline v1
Approach: Pair each input prompt with an expected outcome.
Run the evaluation and review three scores per response:

- Groundedness – did the answer stay inside the knowledge?
- Relevance – did it actually address the question?
- Accuracy – did it match the expected outcome?

Sample evaluation items:

Input: Show all journals with ControlFlag = Fail in period 2026-04
Expect: Returns rows filtered by Period and ControlFlag=Fail, cites JournalIDs.

Input: Which cost centres had a variance over 10% vs budget?
Expect: Returns rows where VarianceVsBudgetPct > 10, ordered desc, with NetUSD cited.

Input: Flag journal J-2271 for audit review – undocumented vendor accrual
Expect: Triggers the Investigate Control Exception topic and calls Flag Journal for Audit Review.

9 Test in the Test Pane

Use the Test pane (right side of the canvas) while you build. For each query, click Activity map to inspect what the agent did:

- Which topic, tool, or knowledge source fired (and in what order)
- Inputs the agent collected and outputs the tool returned
- Citations linking each answer back to a row of your knowledge file
- Latency per step — useful when tuning multi-tool flows

10 Publish to a Channel

When you're happy, click Publish (top right). Then open Channels and enable the surface your users live in. Recommended for this build: Microsoft Teams (Finance Controls channel) and M365 Copilot.

- Microsoft Teams — pin in the Finance Controls channel for auditors
- Microsoft 365 Copilot — for cost centre owners reviewing their own journals
- Custom website — embed on the Finance intranet 'audit corner' page
- Power BI — link the agent from the period close dashboard

Teams and M365 Copilot channels require tenant admin approval before users see the agent — plan for that lead time.

Try These Live

Part 5 of 5 — Sample queries to test your agent during the build-along.

Control Exceptions

Show all journals flagged Fail in the current period.
Which cost centres have the most Review flags this period?

Variance Lookups

What is the largest debit posted to cost centre CC-440 this period?
Show journals with variance > 10% vs budget, biggest absolute value first.

Governed Action

Flag journal J-2271 for review – auditor note 'evidence missing'.
Try a paraphrase: 'this entry looks off, raise it with audit'.

How to Think About This Agent

This agent replaces:

- Manual pivots over the GL extract
- Email chains to chase exception explanation
- Audit findings logged in unstructured spreadsheet

With:

- ✓ Citations linking every number to a JournalID
- ✓ Structured audit findings flowing into SharePoint
- ✓ Faster close, with control exceptions surfaced first

Trusted, governed, enterprise AI — built once, used everywhere.

Tip for real-world use

Bind the agent to your authenticated identity via Entra so a cost-centre owner only sees their own rows automatically. Run the evaluation set after every period close to confirm the agent still answers the top-20 control questions exactly the same way.

You've built a Copilot Studio agent ✓

You now have an AI agent that can:

- Cite every number to a journal row
- Flag control exceptions formally
- Surface variance and control risk first
- Speed up the audit trail